

Module 3 - Data Limitations

Abdullah Sarwar | EN.605.256.82.FA26

The GradCafe data are a self-selected convenience sample rather than a random sample of graduate applicants. People choose whether to post their results, and posting behavior may differ by institution, program, applicant background, or admission outcome. The entries are also self-reported and may contain incorrect, incomplete, or inconsistent information. As a result, counts such as the number of Fall 2026 applications and percentages such as the observed acceptance share describe only the GradCafe records collected for this project. They should not be interpreted as official institutional admission rates or as representative of the full population of graduate applicants.

Missing data also limit several analyses because GPA, GRE, nationality, and some other fields are not available for every record, and available values may use inconsistent wording or formats. The average-score queries therefore exclude SQL NULL values independently for each metric, while Q10 restricts its denominator to records with a usable status; these choices make the calculations reproducible but do not eliminate possible missing-data bias. Q11 depends on LLM-standardized university names, and the four records added during the final Pull Data run do not yet have standardized program or university values, so those rows do not contribute to that grouping. These limitations mean

the database is useful for describing patterns in the collected GradCafe records, but the results should not be generalized beyond this dataset without additional validation and a more representative sampling process.